



Bash basics

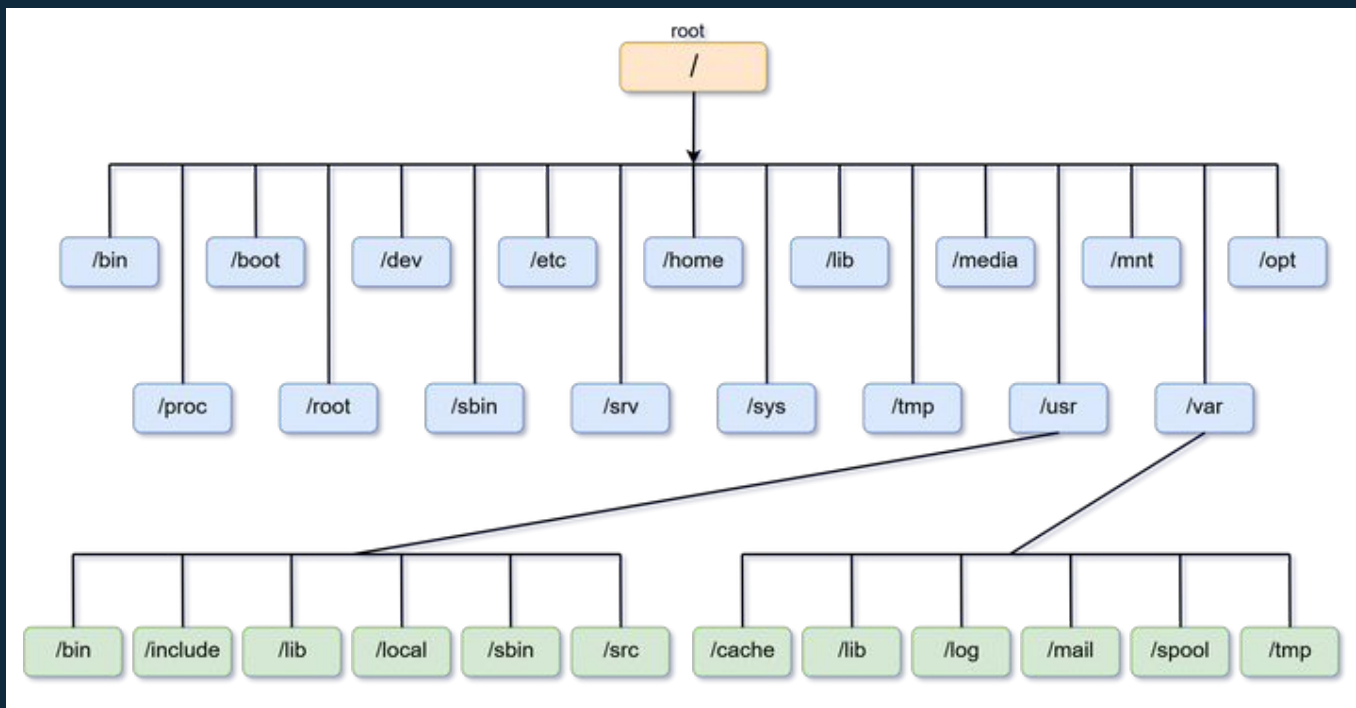


A decorative pattern of hexagons in various shades of blue and cyan on the left side of the slide. Some hexagons contain icons: a lightbulb, a thumbs up, a smartphone, a magnifying glass, a gear, and a speech bubble. A large cyan hexagon in the center of this pattern contains the number '0'.

0

Filesystem Structure

Know it like the back of your hand :)





Directories

/bin /usr/bin

Common executable programs & commands

/usr/sbin

Also holds commands but only admin user specific commands

/home

User home directories
Each user has its own directory

/lib

Shared libraries and kernel modules necessary to boot the system

/proc

Virtual filesystem documenting kernel and process status as text files

/boot

Stores startup files and kernel



Directories

/etc

Configuration files for your system and services

/opt

Optional

Contains extra and third-party software

/var

Variable

Stores files that the system writes and reads data from during operation

/dev

Devices

Stores files that represent the physical parts of the computer

/mnt

Mount

Used for temporarily mounting things

/media

Accessing removable media



File components

- ◇ Everything in linux is a file, including directories (we'll explain later)
- ◇ What defines a file?
 - Name
 - Contents
 - Administrative information
 - Stored in the inode
 - Inodes r wack... we'll come back to them later



A decorative pattern of hexagons in various shades of blue and cyan. Some hexagons contain icons: a lightbulb, a thumbs up, a network of nodes, a smartphone, a magnifying glass, a gear, and a speech bubble. The number '1' is centered in a large cyan hexagon.

1

Review Basic Cmds

Let's start with the basics!



Basic Cmd Navigation



- ◇ pwd
 - Print working directory
- ◇ ls
 - Listing files
- ◇ cd
 - Change directory
 - Absolute Paths vs Relative Paths





File Manipulation



- ◇ touch [FILE]
 - Opens and closes a file
- ◇ cp [SOURCE] [DESTINATION]
- ◇ mv [SOURCE] [DESTINATION]
 - Renaming tool
- ◇ rm [FILE]
 - remove





Directories

- ◇ mkdir [DIRECTORY]
- ◇ rmdir [DIRECTORY]
 - Only works if dir is empty





File Editing



- ◇ gedit [FILE]
 - Common graphical editor
- ◇ nano [FILE]
 - Terminal editor





File viewing

- ◇ `cat [FILE]`
 - Outputs file to stdout
- ◇ `less [FILE]`
 - Scroll thru starting at the top
 - Quit with q
- ◇ `head [FILE]`
 - First 10 lines of file
- ◇ `tail [FILE]`
 - Last 10 lines of file





Input & Output Streams

Standard Input

- ◇ Characters you type into the terminal
- ◇ Input data

Standard Output

- ◇ Regular output printed to the terminal

Standard Error

- ◇ Error output printed to the terminal



- <
 - Redirect standard input to read from a file
- >
 - Redirect standard output to print to a file
 - Overwrites file if exists
- >>
 - Redirect standard output to print to a file
 - Appends to file if exists
- |
 - Pipe operator
 - Redirect standard output to go to another command as standard input



Input & Output Streams

Standard Input

```
sort < [FILE]
```

- ◆ Standard input of sort is taken from the given file

Standard Output

```
echo "Troy Cyber!" >  
testfile
```

- ◆ Standard output redirected to "testfile"

Pipe Operator

```
grep -R "Cyber" testfile | wc -l
```

- ◆ Count the occurrences of "Cyber" in the file "testfile"
- ◆ Output from the left becomes input into the right

A decorative graphic on the left side of the slide. It features a large cyan hexagon in the center containing the number '2'. Surrounding this central hexagon are several smaller hexagons of varying shades of blue and cyan. Some of these smaller hexagons contain white icons: a lightbulb, a thumbs-up, a smartphone, a magnifying glass, and a gear. There is also a network-like icon with a central node and several smaller nodes connected by lines.

2

General Unix Tools

Familiarize yourself with these!



Comparing Files

diff [FILE1] [FILE2]

- ◇ > indicates the line is in file2, but not file 1
- ◇ < indicates the line is in file1, but not file2





Downloading files

wget [URL]

- ◇ Will retrieve whatever file is stored at that url





Extracting files

- ◇ .tar.gz files:
 - `tar -xzf [FILE.tar.gz]`
- ◇ .zip file:
 - `unzip [FILE.zip]`





Searching

- ◇ find
 - Recursively searches based on certain criteria
- ◇ locate
 - uses an indexed database
 - must use updatedb first
- ◇ grep
 - pattern matching throughout text





Misc. useful cmds (good for forensics!)



wc [FILE]

- ◇ Word count

wc -l /etc/passwd

- ◇ Number of users in the passwd file

cut [OPTION] [FILE]

- ◇ Cuts out sections and writes result to standard output

eog [IMAGE]

- ◇ Image viewer

mpg123 [MP3 FILE]

- ◇ Audio player

evince [FILE]

- ◇ Document viewer

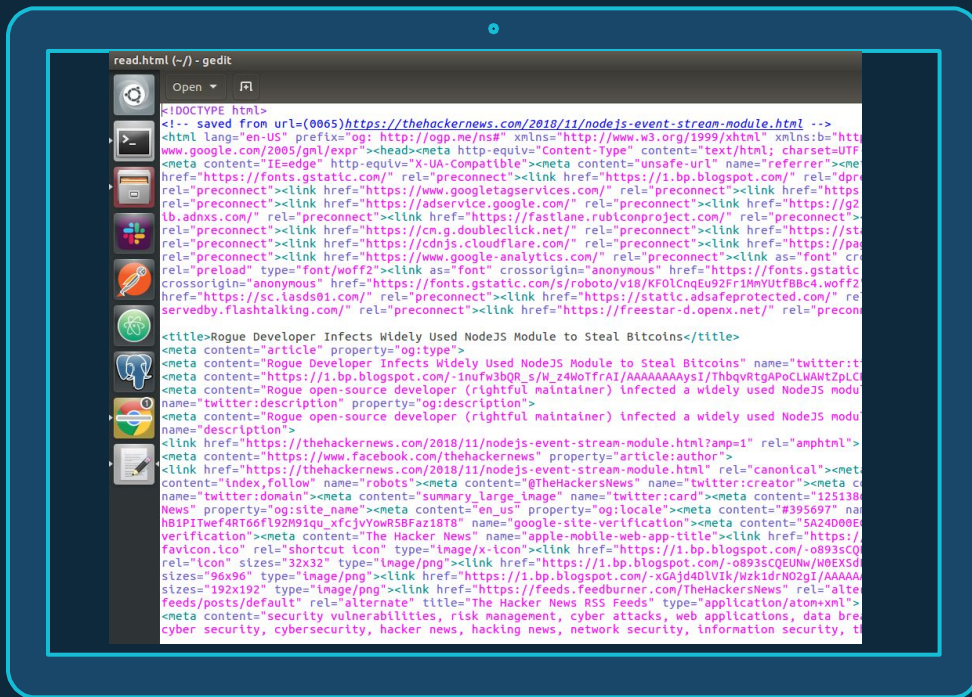


xdg-open {FILE | URL}

- ◆ Yeah you don't need eog, evince, or mpg123 because xdg-open does it ALL
- ◆ Open file or URL in preferred application
- ◆ Supports ftp,file,https,http URLs

Example:

- ◆ xdg-open read.html





Looking up any other commands

- ◇ `man [CMD]`
 - Look through the built in linux manual
- ◇ `whatis [CMD]`
 - very brief description
- ◇ Online man pages
 - Manuals but now with colors and ctrl f functionality :D

MAN



A decorative graphic on the left side of the slide. It features a large cyan hexagon in the center with the number '3' inside. Surrounding this central hexagon are several smaller hexagons of varying shades of blue and cyan. Some of these smaller hexagons contain white icons: a lightbulb, a thumbs-up, a smartphone, a magnifying glass, and a gear. There is also a network-like icon with a central node and radiating lines, and a speech bubble icon. The background of the entire slide is a dark navy blue.

3

Shell Stuff



Bash

- ◇ Bourne-again shell
 - This is the shell you use by default in linux
- ◇ Can be accessed with the bash cmd
- ◇ You can use other shells such as sh if you want
 - Not recommended
- ◇ Bash is a language, cmds are just part of the it





\$	Shell Variable
~	Special home directory variable
&	Background command execution
;	Command termination
* ? []	Shell wildcards
' " \	Metacharacter quotes

Shell Metacharacters


Special Characters used to represent something

- ◇ Redirection symbols we talked about earlier
- ◇ Wildcard substitutions
- ◇ Escape characters





Examples

- ◇ `cd ~`
 - will take you to your user's home directory
 - ◇ `apt upgrade &`
 - will upgrade in the background
 - ◇ `apt update; apt upgrade`
 - will run `apt update.` ; signifies that the command has ended, and then linux will move on to the next one: `apt upgrade`
 - ◇ `ls *.txt`
 - ◇ `ls -l file[1-3]`
 - will list out file1, file2, and file3
 - ◇ `ls -l file?`
 - will list out anything starting w file and having 1 extra character afterwards
 - ◇ `ls file\ 1`
 - will list a file called "file 1". the `\[space]` is the escape character
- 



Shell Metacharacters con't

&&	and operator
	or operator
!	not operator
!!	previous cmd
{ }	create range





Examples

- ◇ Ex: `apt update && apt upgrade -y`
 - Will execute `apt update`. If that works, move onto `apt upgrade -y`. If that works, return true
 - If EITHER of those don't work, return false
- ◇ Ex: `apt-update || apt upgrade`
 - Will execute `apt update` and then `apt upgrade`. If either works, then return true
 - If BOTH don't work, then return false
- ◇ `!` is used more for bash scripting, we'll get there eventually
- ◇ `sudo !!` means run prev cmd as sudo
- ◇ `touch {a..z}` creates files a-z





Shell Variables

- ◇ 2 types
 - Local & environment

env or printenv

- ◇ List of all environmental variables
- ◇ External command → runs in child process

set

- ◇ Display ALL the variables available in the current shell
- ◇ Built-in command





Environmental Variables

- ◇ Defined for current shell & inherited by any child shells or processes
- ◇ `echo $LOGNAME`
 - Display username
- ◇ `echo $HOME`
 - Current user's home directory
- ◇ `echo $PATH`
 - List of directories to search for executable files when the user runs a command





Path variable


- ◇ Ex:
/home/joseph/local/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/usr/local/games
- ◇ When I use ls, the system checks \$PATH and asks, “is ls in any of these directories?”
 - Happens to be in /bin/ls
- ◇ So the system will do /bin/ls
- ◇ Otherwise, the system will check within your current directory for an executable called ls
- ◇ If you wanted to, you could just use /bin/ls directly



A decorative graphic on the left side of the slide. It features a large cyan hexagon in the center containing the number '4'. Surrounding this central hexagon are several smaller hexagons of varying shades of blue and cyan. Some of these smaller hexagons contain white icons: a lightbulb, a thumbs-up, a smartphone, a magnifying glass, and a gear. There is also a network-like icon with a central node and radiating lines. The background of the entire slide is a dark blue gradient.

4


Files in depth : inodes



What is in an inode?

- ◇ Contains administrative info/system data
 - Mode/permission (protection)
 - Owner ID
 - Group ID
 - Size of file
 - Number of hard links to the file (we'll get to this later)
 - Disk block location of file contents
 - Time FILE last accessed
 - Time FILE last modified
 - Time INODE last modified
- ◇ Does NOT contain the FILENAME
 - That's stored in the file's PARENT directory





Why bother understanding inodes?

- ◇ Directory hierarchy is just a convenient way to NAME the files
- ◇ The system's internal name for a file is the i-number
 - I-number: Number of the inode holding the file's info
 - Basically, from the i-number, the kernel can access inode contents (including location of the file) and THEN it can access the file
 - View i-number with `ls -li`
- ◇ Could be potentially tested as a forensics question
- ◇ Are one of the foundational elements of the linux FS
- ◇ Responsible for hardlinks in linux





Example

```
joseph@pop-os:~/Desktop$ ls -li
total 428
3014734 drwxr-xr-x 5 joseph joseph 4096 Aug 23 11:05 ctf
2491798 drwxrwxr-x 7 joseph joseph 4096 Jul 24 22:53 internship
413101 -rw-rw-r-- 1 joseph joseph 243100 Aug 23 17:25 'Joseph Xu Resume
f'
407746 -rw-r--r-- 1 joseph joseph 0 Sep 14 10:47 junk
3021400 drwxr-xr-x 6 joseph joseph 4096 Aug 13 17:34 RsaCtfTool
789575 drwxr-xr-x 8 joseph joseph 4096 Aug 14 22:58 textbooks
424325 -rw----- 1 joseph joseph 167981 Aug 23 18:21 Transcripts.pdf
```

- ◆ Number on the very left is the inode-number
 - This is how the system refers to files in linux





So what's in a directory?

- ◇ Directories are just files that contain special tables
 - Contain filenames within the directory and their corresponding inode numbers
 - So a filename in a directory is called a LINK because it links a name in the directory to the inode
 - First 2 entries are always . and ..
 - . = inode of CURRENT directory
 - .. = inode of PARENT directory
- ◇ Unfortunately, linux does not allow us to view the raw contents of a directory w/o special tools





Weird implications of how directories work

- ◇ Same inode can appear more than once in a directory, with 2 separate links
- ◇ `rm` command does not remove inodes
 - Removes the links (directory entries)
 - System only removes inode when LAST link to a file disappears





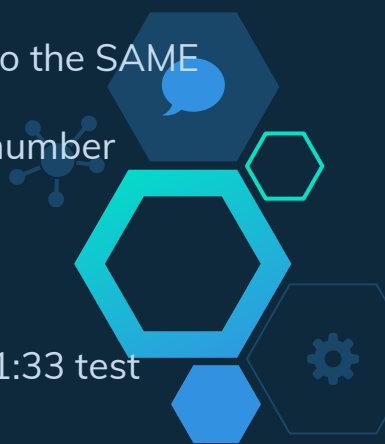
Analogy using Java

- ◇ In Java, we have objects
 - The content of an object is analogous to file contents
- ◇ Java compiler refers to objects using memory addresses
 - Compiler = linux kernel
 - Memory address = inode
- ◇ WE refer to objects using variables, which just contain memory addresses
 - Variables = link (filename)





Application of inodes

- ◇ Hard links
 - ◇ File contents can have MULTIPLE names (aka links)
 - When we do this, we are creating hardlinks
 - Ex: /home/joseph/file1 and /home/joseph/file2 can refer to the SAME file contents
 - Because the filename can link to the SAME inode number
 - ◇ Create using ln command
 - ln [original] [link]
 - ◇ View # of hardlinks using ls -l
 - Comes right after the permissions
 - Ex: 3407989 drwxr-xr-x 2 joseph joseph 4096 Sep 14 11:33 test
- 



Java analogy part 2

- ◇ So if you have:
 - `Object var1 = new Object();`
 - `Object var2 = var1;`
- ◇ Var2 and var1 stores the same memory address
 - Therefore, the object we created has 2 “filenames” (links): var1 and var2
 - When var1 is modified, so is var2
 - This is kinda important for all of y’all struggling with AP CS A





Linking for directories

- ◇ They work a bit differently
 - Can't use ln with directories
- ◇ Each directory starts out with 2 links
 - Itself (. directory)
 - One from the parent directory
- ◇ Each subdirectory counts as another additional link
 - Because each subdirectory has a .. entry
- ◇ So $2 + \text{\#subdirectories} = \text{total amount of links to a directory}$



```

joseph@pop-os:~/Desktop/inodes$ mkdir test
joseph@pop-os:~/Desktop/inodes$ ls -lai
total 12
3407987 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:36 .
393300 drwxr-xr-x 8 joseph joseph 4096 Sep 14 11:36 ..
3407989 drwxr-xr-x 2 joseph joseph 4096 Sep 14 11:36 test
joseph@pop-os:~/Desktop/inodes$ touch file1; echo "hello" > file1
joseph@pop-os:~/Desktop/inodes$ ls -lai
total 16
3407987 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:36 .
393300 drwxr-xr-x 8 joseph joseph 4096 Sep 14 11:36 ..
3407990 -rw-r--r-- 1 joseph joseph 6 Sep 14 11:36 file1
3407989 drwxr-xr-x 2 joseph joseph 4096 Sep 14 11:36 test
joseph@pop-os:~/Desktop/inodes$ ln file1 file2
joseph@pop-os:~/Desktop/inodes$ head file2
hello
joseph@pop-os:~/Desktop/inodes$ ls -lai
total 20
3407987 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:36 .
393300 drwxr-xr-x 8 joseph joseph 4096 Sep 14 11:36 ..
3407990 -rw-r--r-- 2 joseph joseph 6 Sep 14 11:36 file1
3407990 -rw-r--r-- 2 joseph joseph 6 Sep 14 11:36 file2
3407989 drwxr-xr-x 2 joseph joseph 4096 Sep 14 11:36 test
joseph@pop-os:~/Desktop/inodes$ mkdir test/reference
joseph@pop-os:~/Desktop/inodes$ ls -lai
total 20
3407987 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:36 .
393300 drwxr-xr-x 8 joseph joseph 4096 Sep 14 11:36 ..
3407990 -rw-r--r-- 2 joseph joseph 6 Sep 14 11:36 file1
3407990 -rw-r--r-- 2 joseph joseph 6 Sep 14 11:36 file2
3407989 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:37 test
joseph@pop-os:~/Desktop/inodes$ ls -ldi
3407987 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:36 .
joseph@pop-os:~/Desktop/inodes$

```

- ◇ Created test, which has 2 links
- ◇ created file1 with “hello” inside
 - Currently has 1 link, which is just itself
- ◇ Created hardlink file2, which points to file1’s inode
 - When we view file2 contents, it’s the same as file1
- ◇ When we view the inode/hardlink info with ls -lai
 - file1 and file2 have the same inode number: 3407990
 - 3407990 has 2 links to it: file1 and file2
- ◇ When we create test/reference, test’s hardlink count raises by 1 because reference contains the entry .., which links to test’s inode





Implications of hardlinks

- ◇ Ex: file1 and file2 were linked to the same inode
 - When you change file1, file2 changes as well
 - And vice versa
 - When you remove file1, the original, we still have file2 remaining and we can still view the file contents
- ◇ cp in linux creates a whole NEW file, with a different inode and just copies over the content from the original





Security implications of hardlinks

- ◇ Common issue: unauthorized file
 - But if that unauthorized file had a hardlink to it, then the content remains on the system, just in a different location





Symlinks

- ◇ Similar to hardlinks in that it will refer to an already existing file
 - But VERY different implementation
 - Does NOT have the same inode
- ◇ It's just a pointer that redirects to a different file path



Example

```
joseph@pop-os:~/Desktop/inodes$ ln -s file1 file3
joseph@pop-os:~/Desktop/inodes$ ls -lai
total 20
3407987 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:53 .
393300 drwxr-xr-x 8 joseph joseph 4096 Sep 14 11:36 ..
3407990 -rw-r--r-- 2 joseph joseph 6 Sep 14 11:36 file1
3407990 -rw-r--r-- 2 joseph joseph 6 Sep 14 11:36 file2
3407992 lrwxrwxrwx 1 joseph joseph 5 Sep 14 11:53 file3 -> file1
3407989 drwxr-xr-x 3 joseph joseph 4096 Sep 14 11:37 test
joseph@pop-os:~/Desktop/inodes$ head file3
hello
```

- ◇ -s means symbolic link
- ◇ file3 -> file1
 - Arrow indicates symlink
 - file1 has inode 3407990
 - file3 has inode 3407992
- ◇ When you view file3, it redirects you to file1, and therefore you see the file1 contents



Implications of symlinks

- ◇ file3 is symlinked to file1
- ◇ When you modify file1 or file3, you are modifying the SAME content
- ◇ When you remove file3, file1 remains and you can use it normally
- ◇ BUT when you remove file1, file3 still tries to redirect you to file1
 - Dangling symlink





Security implications of symlinks

- ◇ If you find a backdoor's SYMLINK and you remove it, then the backdoor remains on your system
 - Find where the ACTUAL location is with `ls -l` and look for where the symlink points to





Thanks!

Any questions?

